

Salty Embeddings

Privacy Preserving RAG on MongoDB Atlas


Vitaly Kleban, San-Francisco, 2024
averagejoe.ai

salt prevents brute-force attacks

$\text{sha512}(\text{password}) \rightarrow \text{sha512}(\text{password} + \text{salt})$

embeddings \rightarrow ???

Source



Reconstructed



client side encryption prevents data leakage
queryable encryption makes search efficient

$\{\text{data}, \text{embedding}\} \rightarrow \{\text{encrypt}(\text{data}), \text{embedding}\}$

```
# Query embedding
query_emb = np.random.rand(128)

# Document embedding
embedding = np.random.rand(128)

# Salt - fixed random permutation of embedding elements
salt = np.random.permutation(len(embedding))

salty_query = query_emb[salt]
salty_embedding = embedding[salt]

# Check whether distance is preserved
query_emb.dot(embedding), salty_query.dot(salty_embedding)
```

✓ 0.0s

Python

(35.446370272374715, 35.446370272374715)

“Standard” encryption destroys distance
between vectors and prevents vector search

Our options:
homomorphic encryption,
secure multiparty computation,
locality-sensitive hashing,
functional encryption,
random projection, permutation

Client side encryption 

+

Queryable encryption 

+

Salty embeddings 

=

Privacy Preserving RAG 